

```
In [105]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import scipy.stats as stats
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn import metrics
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.preprocessing import scale
from sklearn.metrics import roc_auc_score, roc_curve, classification_report, c
```

```
In [106]: df = pd.read_csv("Holiday_Package.csv")
```

```
In [4]: df.head()
```

Out[4]:

|   | Unnamed: 0 | Holliday_Package | Salary | age | educ | no_young_children | no_older_children | foreign |
|---|------------|------------------|--------|-----|------|-------------------|-------------------|---------|
| 0 | 1          | no               | 48412  | 30  | 8    | 1                 | 1                 | no      |
| 1 | 2          | yes              | 37207  | 45  | 8    | 0                 | 1                 | no      |
| 2 | 3          | no               | 58022  | 46  | 9    | 0                 | 0                 | no      |
| 3 | 4          | no               | 66503  | 31  | 11   | 2                 | 0                 | no      |
| 4 | 5          | no               | 66734  | 44  | 12   | 0                 | 2                 | no      |

```
In [5]: #Dropping unnecessary column:
df=df.drop('Unnamed: 0',axis=1)
```

```
In [6]: df.head()
```

Out[6]:

|   | Holliday_Package | Salary | age | educ | no_young_children | no_older_children | foreign |
|---|------------------|--------|-----|------|-------------------|-------------------|---------|
| 0 | no               | 48412  | 30  | 8    | 1                 | 1                 | no      |
| 1 | yes              | 37207  | 45  | 8    | 0                 | 1                 | no      |
| 2 | no               | 58022  | 46  | 9    | 0                 | 0                 | no      |
| 3 | no               | 66503  | 31  | 11   | 2                 | 0                 | no      |
| 4 | no               | 66734  | 44  | 12   | 0                 | 2                 | no      |

```
In [7]: df.shape
```

Out[7]: (872, 7)

```
In [119]: #checking data types:
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 872 entries, 0 to 871
Data columns (total 8 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Unnamed: 0            872 non-null    int64
 1   Holliday_Package      872 non-null    object
 2   Salary                872 non-null    int64
 3   age                  872 non-null    int64
 4   educ                 872 non-null    int64
 5   no_young_children    872 non-null    int64
 6   no_older_children    872 non-null    int64
 7   foreign              872 non-null    object
dtypes: int64(6), object(2)
memory usage: 54.6+ KB
```

```
In [9]: #Checking for missing values:
df.isnull().sum()
```

```
Out[9]: Holliday_Package    0
Salary                    0
age                      0
educ                     0
no_young_children        0
no_older_children        0
foreign                  0
dtype: int64
```

```
In [10]: df.isna().sum()
```

```
Out[10]: Holliday_Package    0
Salary                    0
age                      0
educ                     0
no_young_children        0
no_older_children        0
foreign                  0
dtype: int64
```

```
In [11]: #Checking for duplicate values:
df.duplicated().sum()
```

```
Out[11]: 0
```

```
In [12]: #Checking statistical summary:
df.describe(include='all')
```

Out[12]:

|        | Holliday_Package | Salary        | age        | educ       | no_young_children | no_older_child |
|--------|------------------|---------------|------------|------------|-------------------|----------------|
| count  | 872              | 872.000000    | 872.000000 | 872.000000 | 872.000000        | 872.000000     |
| unique | 2                | NaN           | NaN        | NaN        | NaN               | NaN            |
| top    | no               | NaN           | NaN        | NaN        | NaN               | NaN            |
| freq   | 471              | NaN           | NaN        | NaN        | NaN               | NaN            |
| mean   | NaN              | 47729.172018  | 39.955275  | 9.307339   | 0.311927          | 0.982706       |
| std    | NaN              | 23418.668531  | 10.551675  | 3.036259   | 0.612870          | 1.086430       |
| min    | NaN              | 1322.000000   | 20.000000  | 1.000000   | 0.000000          | 0.000000       |
| 25%    | NaN              | 35324.000000  | 32.000000  | 8.000000   | 0.000000          | 0.000000       |
| 50%    | NaN              | 41903.500000  | 39.000000  | 9.000000   | 0.000000          | 1.000000       |
| 75%    | NaN              | 53469.500000  | 48.000000  | 12.000000  | 0.000000          | 2.000000       |
| max    | NaN              | 236961.000000 | 62.000000  | 21.000000  | 3.000000          | 6.000000       |

```
In [13]: #Checking unique counts for all variables:
for feature in df.columns:
    print(feature)
    print(df[feature].value_counts())
    print(df[feature].nunique())
    print('\n')
```

```
Holliday_Package
no      471
yes     401
Name: Holliday_Package, dtype: int64
2
```


```
Salary
46195      2
33357      2
39460      2
36976      2
40270      2
..
38352      1
119644     1
96072      1
115431     1
74659      1
Name: Salary, dtype: int64
```

```
In [14]: #Checking proportion of target variable:  
df.Holliday_Package.value_counts(normalize=True)
```

```
Out[14]: no      0.540138  
yes       0.459862  
Name: Holliday_Package, dtype: float64
```

```
In [ ]:
```

```
In [16]: #Checking correlation between the continuous variables:  
plt.figure(figsize=(10,8))  
sns.set(font_scale=1.2)  
sns.heatmap(df.corr(),annot=True,)  
plt.show()
```

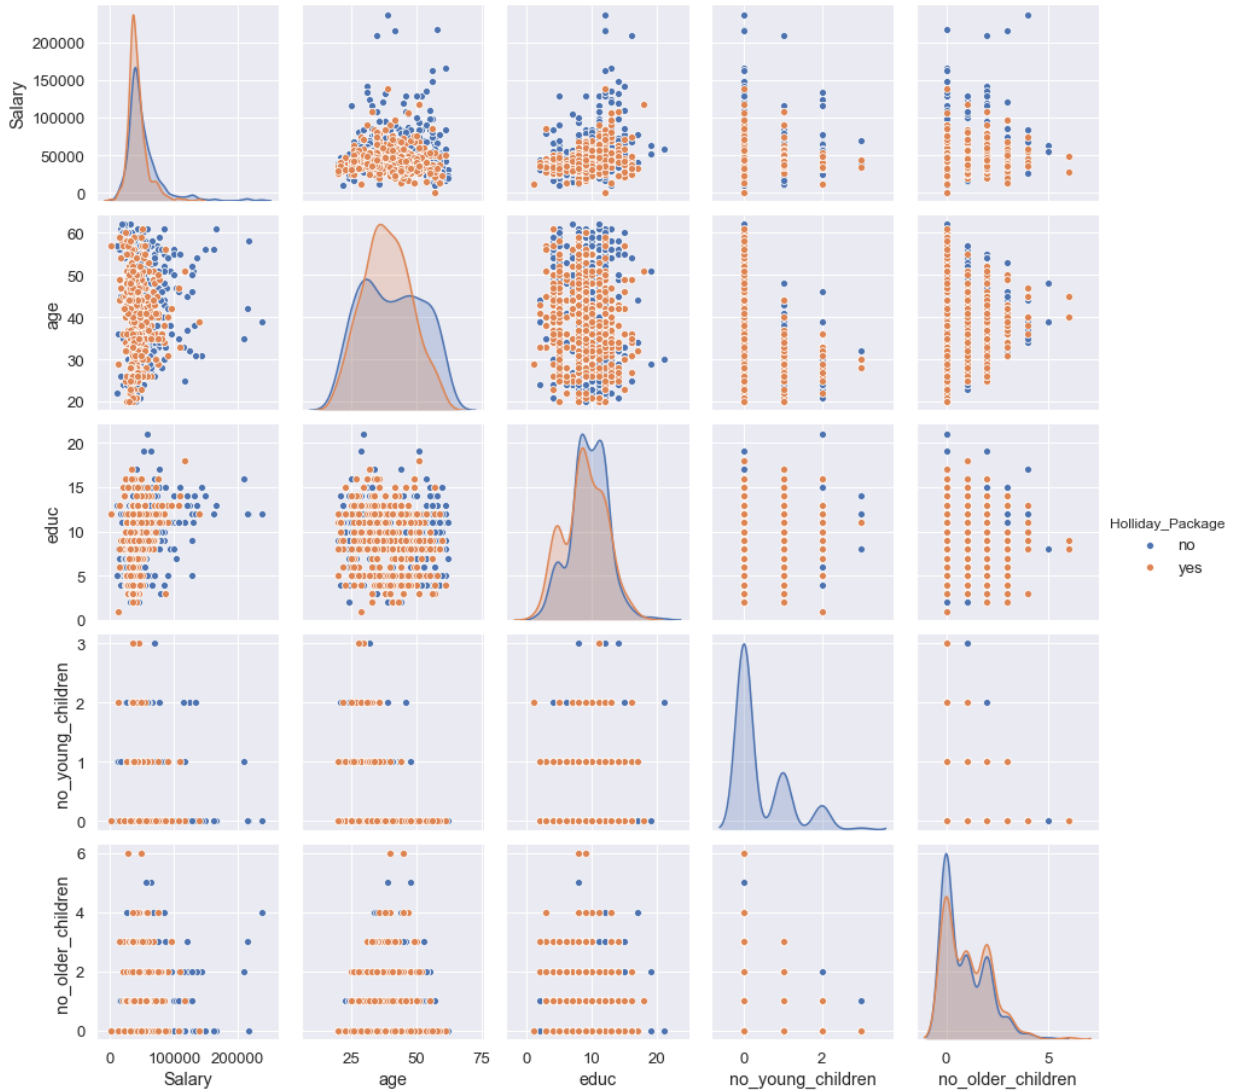


```
In [69]: sns.pairplot(df, hue='Holliday_Package', diag_kind='kde')
```

```
/opt/anaconda3/anaconda3/lib/python3.8/site-packages/seaborn/distribution
s.py:369: UserWarning: Default bandwidth for data is 0; skipping density
estimation.
```

```
warnings.warn(msg, UserWarning)
```

```
Out[69]: <seaborn.axisgrid.PairGrid at 0x125b56be0>
```




```
In [17]: pd.crosstab(df.foreign,df.Holliday_Package).plot(kind='bar')
plt.title('Foreign Vs. Holiday Package')
plt.xlabel('Foreign')
plt.ylabel('Holiday Package')
```

```
Out[17]: Text(0, 0.5, 'Holiday Package')
```

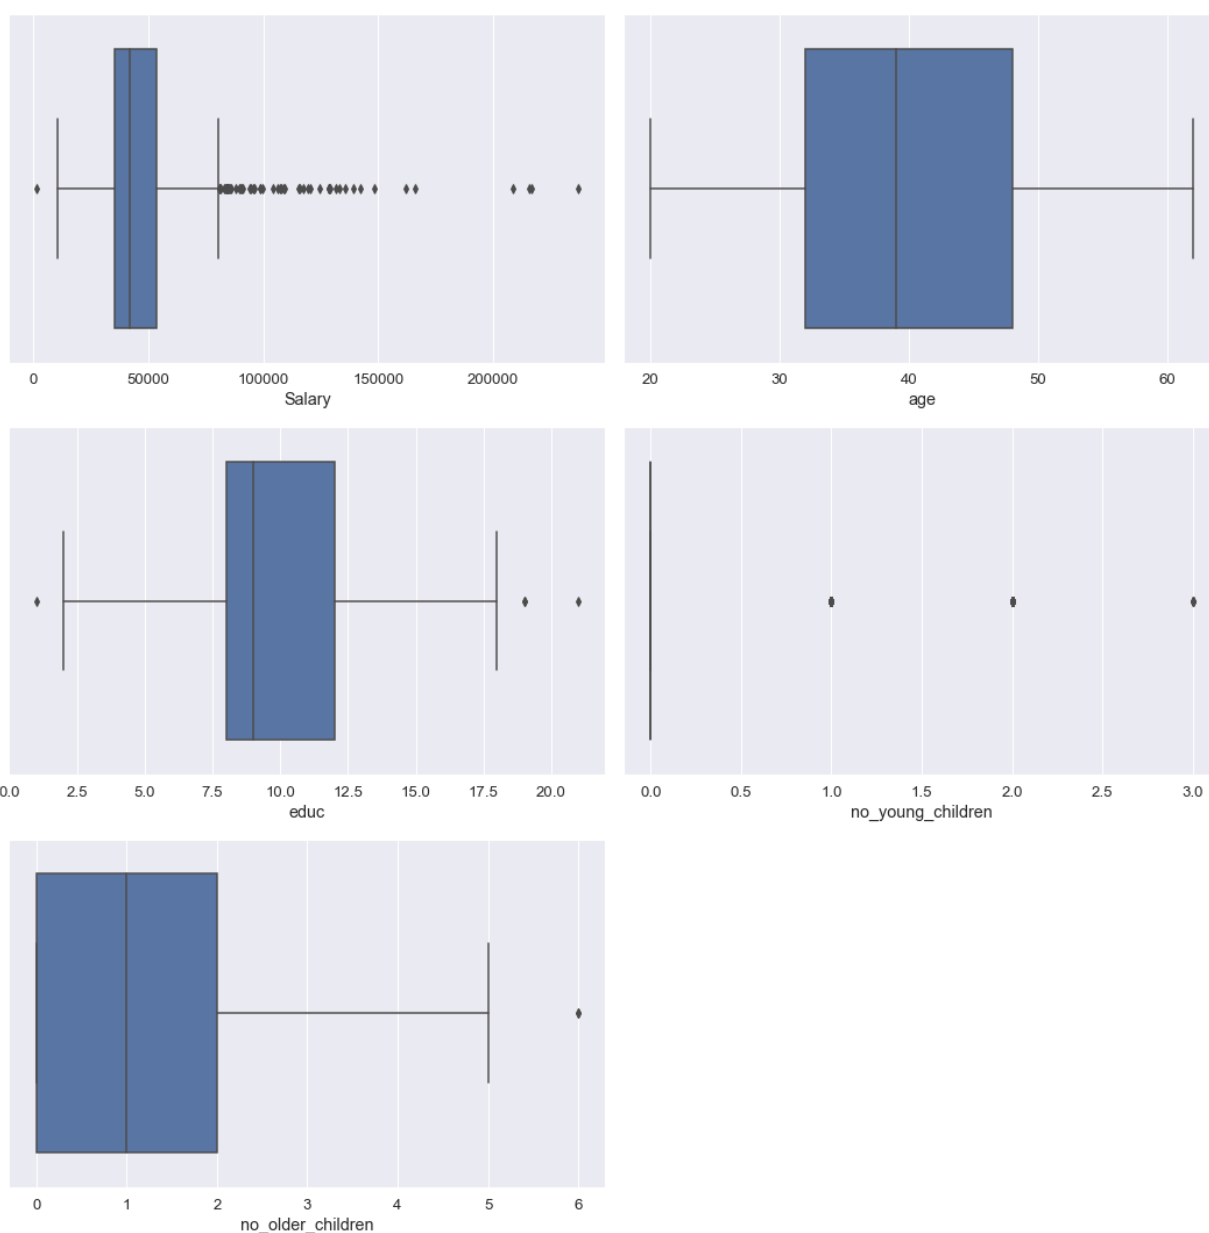


```
In [37]: from pylab import rcParams
rcParams['figure.figsize'] = 20,20
df[['Salary', 'age', 'educ', 'no_young_children', 'no_older_children']].hist();
```



In [51]: *#Checking for outliers in continuous variables:*

```
data_plot=df[['Salary','age','educ','no_young_children','no_older_children']
fig=plt.figure(figsize=(15,15))
for i in range(0,len(data_plot.columns)):
    ax=fig.add_subplot(3,2,i+1)
    sns.boxplot(data_plot[data_plot.columns[i]])
    plt.tight_layout()
```





In [ ]:

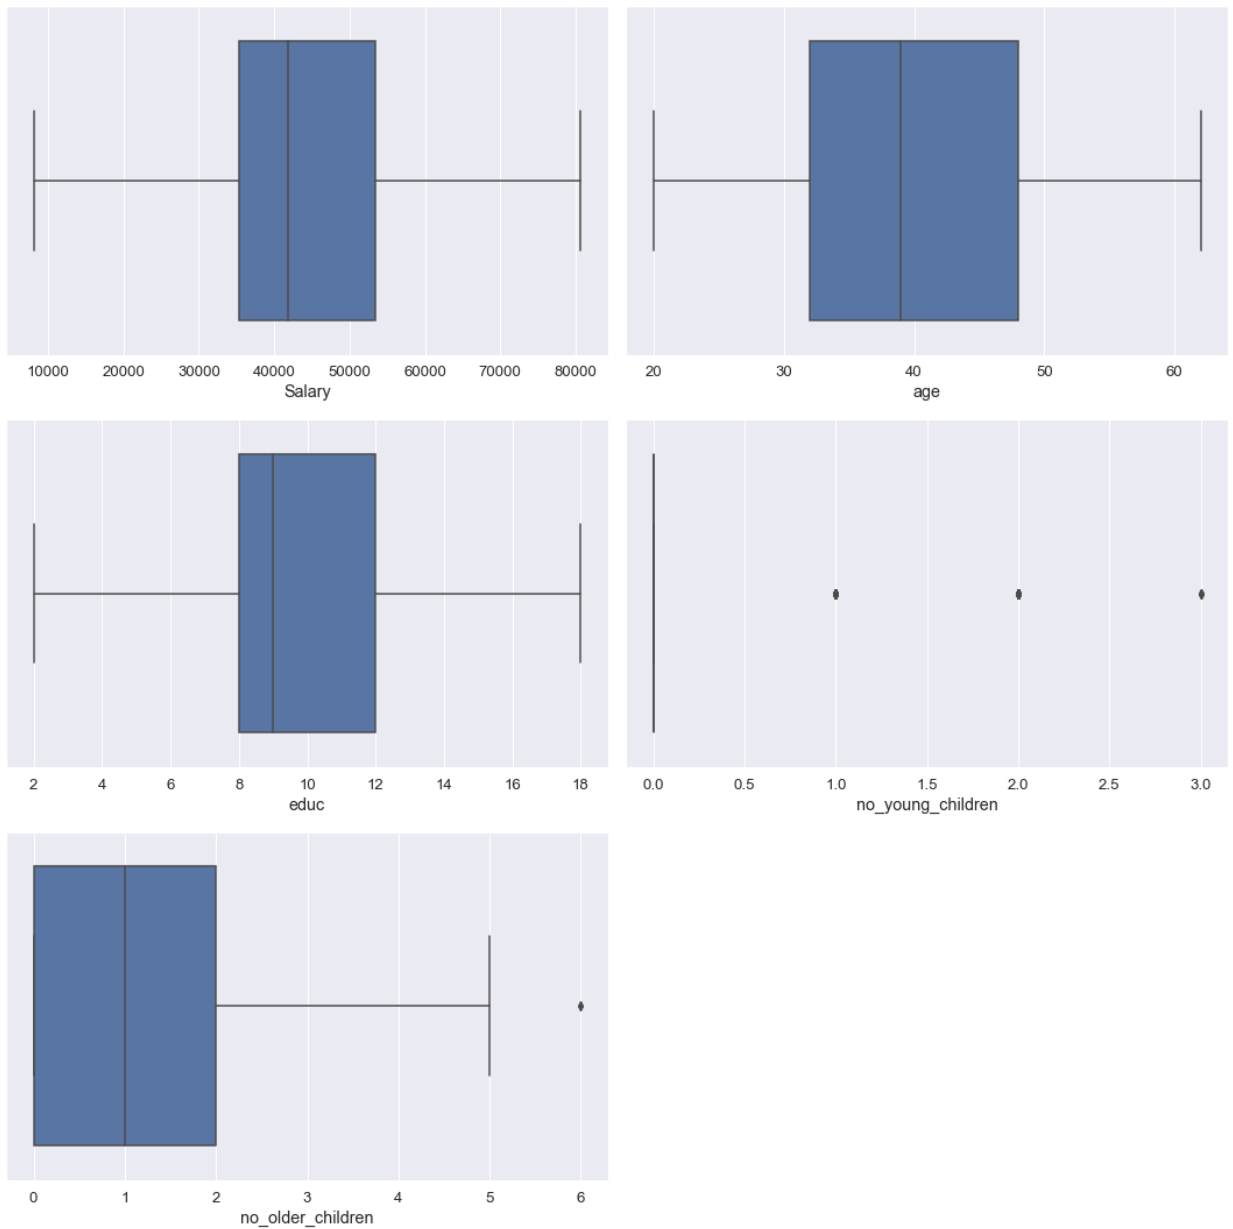
```
In [120]: #Treating the outliers:
def remove_outlier(col):
    sorted(col)
    Q1,Q3=col.quantile([0.25,0.75])
    IQR=Q3-Q1
    lower_range=Q1-(1.5*IQR)
    upper_range=Q3+(1.5*IQR)
    return lower_range,upper_range
```

```
In [121]: #Replacing outliers using capping and flooring

lrsalary,upsalary=remove_outlier(df['Salary'])
df['Salary']=np.where(df['Salary']>upsalary,upsalary,df['Salary'])
df['Salary']=np.where(df['Salary']<lrsalary,lrsalary,df['Salary'])

lredu,upedu=remove_outlier(df['educ'])
df['educ']=np.where(df['educ']>upedu,upedu,df['educ'])
df['educ']=np.where(df['educ']<lredu,lredu,df['educ'])
```

```
In [57]: #Rechecking for outliers:
data_plot=df[['Salary','age','educ','no_young_children','no_older_children']
fig=plt.figure(figsize=(15,15))
for i in range(0,len(data_plot.columns)):
    ax=fig.add_subplot(3,2,i+1)
    sns.boxplot(data_plot[data_plot.columns[i]])
    plt.tight_layout()
```



```
In [122]: #Converting object type data to categorical codes:
code1={"no":0, "yes":1}
df["Holliday_Package"]=df["Holliday_Package"].replace(code1)
```

```
In [123]: code2={"no":0, "yes":1}
df["foreign"]=df["foreign"].replace(code2)
```

```
In [124]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 872 entries, 0 to 871
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            872 non-null   int64
1   Holliday_Package      872 non-null   int64
2   Salary                872 non-null   float64
3   age                   872 non-null   int64
4   educ                  872 non-null   float64
5   no_young_children     872 non-null   int64
6   no_older_children     872 non-null   int64
7   foreign                872 non-null   int64
dtypes: float64(2), int64(6)
memory usage: 54.6 KB
```

```
In [125]: #Cross-checking the categorical coding:
for column in df[['Holliday_Package', 'foreign']]:
    print(df[column].value_counts().sort_values())
    print('\n')
```

```
1    401
0    471
Name: Holliday_Package, dtype: int64
```

```
1    216
0    656
Name: foreign, dtype: int64
```

```
In [63]: #Splitting the data into train and test sets:
x = df.drop('Holliday_Package', axis=1)
y = df[['Holliday_Package']]
```

```
In [65]: #Splitting in the 70:30 ratio:

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.30 ,
```

```
In [66]: # Fitting the Logistic Regression model:
model = LogisticRegression(solver='newton-cg',max_iter=10000,penalty='none')
model.fit(x_train, y_train)
```

```
/opt/anaconda3/anaconda3/lib/python3.8/site-packages/sklearn/utils/validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
```

```
    return f(*args, **kwargs)
[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
```

```
[Parallel(n_jobs=2)]: Done    1 out of    1 | elapsed:    3.6s finished
```

```
Out[66]: LogisticRegression(max_iter=10000, n_jobs=2, penalty='none', solver='newton-cg',
                             verbose=True)
```

```
In [67]: #Predicting on training and test data:
ytrain_predict = model.predict(x_train)
ytest_predict = model.predict(x_test)
```

```
In [68]: #Checking predicted class and probabilities:  
ytest_predict_prob=model.predict_proba(x_test)  
pd.DataFrame(ytest_predict_prob).head()
```

Out[68]:

|   | 0        | 1        |
|---|----------|----------|
| 0 | 0.773162 | 0.226838 |
| 1 | 0.273315 | 0.726685 |
| 2 | 0.903230 | 0.096770 |
| 3 | 0.958335 | 0.041665 |
| 4 | 0.512241 | 0.487759 |

```
In [82]: #Finding accuracy of training data:  
lr_train_acc=model.score(x_train, y_train)  
lr_train_acc
```

Out[82]: 0.6754098360655738

```
In [83]: #Calculating AUC and plotting ROC curve for training data:
probs = model.predict_proba(x_train)
probs = probs[:, 1]


lr_train_auc = roc_auc_score(y_train, probs)
print('AUC: %.3f' % lr_train_auc)

%matplotlib inline
train_fpr, train_tpr, train_thresholds = roc_curve(y_train, probs)
plt.plot([0, 1], [0, 1], linestyle='--')

plt.plot(train_fpr, train_tpr);
plt.title("ROC Curve - Training data")
```

AUC: 0.742

Out[83]: Text(0.5, 1.0, 'ROC Curve - Training data')



```
In [84]: # Finding Accuracy for Test Data:
lr_test_acc=model.score(x_test, y_test)
lr_test_acc
```

Out[84]: 0.6374045801526718

```
In [85]: ##Calculating AUC and plotting ROC curve for test data:
probs = model.predict_proba(x_test)

probs = probs[:, 1]


lr_test_auc = roc_auc_score(y_test, probs)
print('AUC: %.3f' % lr_test_auc)

test_fpr, test_tpr, test_thresholds = roc_curve(y_test, probs)
plt.plot([0, 1], [0, 1], linestyle='--')

plt.plot(test_fpr, test_tpr);
plt.title("ROC Curve - Testing data")
```

AUC: 0.705

Out[85]: Text(0.5, 1.0, 'ROC Curve - Testing data')

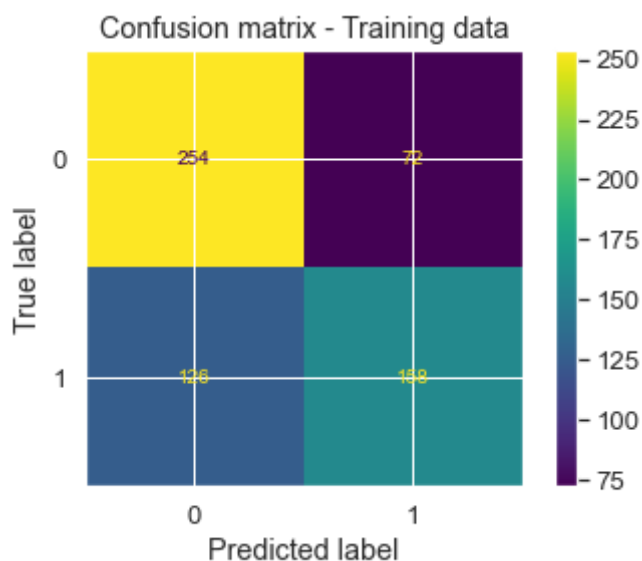


```
In [75]: #Building confusion matrix for training data:
confusion_matrix(y_train, ytrain_predict)
```

Out[75]: array([[252, 74],  
[124, 160]])

```
In [145]: plot_confusion_matrix(model,x_train,y_train);
plt.title("Confusion matrix - Training data")
```

```
Out[145]: Text(0.5, 1.0, 'Confusion matrix - Training data')
```



```
In [77]: print(classification_report(y_train, ytrain_predict))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.67      | 0.77   | 0.72     | 326     |
| 1            | 0.68      | 0.56   | 0.62     | 284     |
| accuracy     |           |        | 0.68     | 610     |
| macro avg    | 0.68      | 0.67   | 0.67     | 610     |
| weighted avg | 0.68      | 0.68   | 0.67     | 610     |

```
In [ ]:
```




```
In [78]: #Building confusion matrix for test data:
confusion_matrix(y_test, ytest_predict)
```

```
Out[78]: array([[102,  43],
               [ 52,  65]])
```

```
In [146]: plot_confusion_matrix(model,x_test,y_test);
plt.title("Confusion matrix - Test data")
```

```
Out[146]: Text(0.5, 1.0, 'Confusion matrix - Test data')
```



```
In [80]: print(classification_report(y_test, ytest_predict))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.66      | 0.70   | 0.68     | 145     |
| 1            | 0.60      | 0.56   | 0.58     | 117     |
| accuracy     |           |        | 0.64     | 262     |
| macro avg    | 0.63      | 0.63   | 0.63     | 262     |
| weighted avg | 0.64      | 0.64   | 0.64     | 262     |

```
In [95]: lr_train_recall=0.56
lr_train_precision=0.68
lr_train_F1=0.62
```

```
In [99]: #Storing data for comparison:
index=['Accuracy', 'AUC', 'Recall','Precision','F1 Score']
data = pd.DataFrame({'Log-reg Train':[lr_train_acc,lr_train_auc,lr_train_re
                                'Log-reg Test':[lr_test_acc,lr_test_auc,'0.56','0.60',
round(data,2)
```

Out[99]:

|                  | Log-reg Train | Log-reg Test |
|------------------|---------------|--------------|
| <b>Accuracy</b>  | 0.68          | 0.637405     |
| <b>AUC</b>       | 0.74          | 0.705158     |
| <b>Recall</b>    | 0.56          | 0.56         |
| <b>Precision</b> | 0.68          | 0.60         |
| <b>F1 Score</b>  | 0.62          | 0.58         |

```
In [127]: #Creating a copy of dataset to perform LDA:
df1=df.copy()
df1.head()
```

Out[127]:

|          | Unnamed: 0 | Holliday_Package | Salary  | age | educ | no_young_children | no_older_children | foreign |
|----------|------------|------------------|---------|-----|------|-------------------|-------------------|---------|
| <b>0</b> | 1          | 0                | 48412.0 | 30  | 8.0  | 1                 | 1                 | 0       |
| <b>1</b> | 2          | 1                | 37207.0 | 45  | 8.0  | 0                 | 1                 | 0       |
| <b>2</b> | 3          | 0                | 58022.0 | 46  | 9.0  | 0                 | 0                 | 0       |
| <b>3</b> | 4          | 0                | 66503.0 | 31  | 11.0 | 2                 | 0                 | 0       |
| <b>4</b> | 5          | 0                | 66734.0 | 44  | 12.0 | 0                 | 2                 | 0       |

```
In [128]: #Dropping unnecessary column:
df1=df1.drop('Unnamed: 0',axis=1)
```

```
In [129]: #Splitting the data into train and test sets:
X = df1.drop('Holliday_Package', axis=1)
Y = df1[['Holliday_Package']]
```

In [113]:

```
In [130]: X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.30,random_
```

```
In [131]: #Building LDA Model:
clf = LinearDiscriminantAnalysis()
model=clf.fit(X_train,Y_train)
```

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/sklearn/utils/validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples, ), for example using ravel().

```
return f(*args, **kwargs)
```

```
In [132]: #Making predictions on the training data with 0.5 cut-off value:
```


```
pred_class_train = model.predict(X_train)
```

```
# Test Data Class Prediction with a cut-off value of 0.5:
pred_class_test = model.predict(X_test)
```

```
In [136]: f,a = plt.subplots(1,2,sharex=True,sharey=True,squeeze=False)
fig=plt.figure(figsize=(20,20))
#Plotting confusion matrix for the different models for the Training Data

plot_0 = sns.heatmap((metrics.confusion_matrix(Y_train,pred_class_train)),a
a[0][0].set_title('Training Data'))

plot_1 = sns.heatmap((metrics.confusion_matrix(Y_test,pred_class_test)),ann
a[0][1].set_title('Test Data'));
```



<Figure size 1440x1440 with 0 Axes>

```
In [137]: #Classification report for training and test:
print('Classification Report of the training data:\n\n',metrics.classification_report(y_train,y_train_pred))
print('Classification Report of the test data:\n\n',metrics.classification_report(y_test,y_test_pred))
```

Classification Report of the training data:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.67      | 0.78   | 0.72     | 326     |
| 1            | 0.69      | 0.56   | 0.61     | 284     |
| accuracy     |           |        | 0.68     | 610     |
| macro avg    | 0.68      | 0.67   | 0.67     | 610     |
| weighted avg | 0.68      | 0.68   | 0.67     | 610     |

Classification Report of the test data:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.66      | 0.70   | 0.68     | 145     |
| 1            | 0.60      | 0.56   | 0.58     | 117     |
| accuracy     |           |        | 0.64     | 262     |
| macro avg    | 0.63      | 0.63   | 0.63     | 262     |
| weighted avg | 0.64      | 0.64   | 0.64     | 262     |

```
In [138]: # Training Data Probability Prediction
pred_prob_train = model.predict_proba(X_train)

# Test Data Probability Prediction
pred_prob_test = model.predict_proba(X_test)
```

```
In [140]: Ytest_predict_prob=model.predict_proba(X_test)
pd.DataFrame(Ytest_predict_prob).head()
```

Out[140]:

|   | 0        | 1        |
|---|----------|----------|
| 0 | 0.763849 | 0.236151 |
| 1 | 0.278109 | 0.721891 |
| 2 | 0.887980 | 0.112020 |
| 3 | 0.950289 | 0.049711 |
| 4 | 0.507535 | 0.492465 |

```
In [141]: # AUC and ROC for the training data
```

```
auc = metrics.roc_auc_score(Y_train,pred_prob_train[:,1])  
print('AUC for the Training Data: %.3f' % auc)
```

```
# calculate and plot roc curve
```

```
fpr, tpr, thresholds = metrics.roc_curve(Y_train,pred_prob_train[:,1])  
plt.plot([0, 1], [0, 1], linestyle='--')  
plt.plot(fpr, tpr, marker='.',label = 'Training Data')
```

```
# AUC and ROC for the test data
```

```
# calculate AUC
```

```
auc = metrics.roc_auc_score(Y_test,pred_prob_test[:,1])  
print('AUC for the Test Data: %.3f' % auc)
```

```
# calculate and plot roc curve
```

```
fpr, tpr, thresholds = metrics.roc_curve(Y_test,pred_prob_test[:,1])  
plt.plot([0, 1], [0, 1], linestyle='--')  
plt.plot(fpr, tpr, marker='.',label='Test Data')
```

```
# show the plot
```

```
plt.legend(loc='best')
```

```
plt.show()
```

AUC for the Training Data: 0.739

AUC for the Test Data: 0.703



In [142]: *#Storing data for comparison:*

```
LDA_train_acc=0.68
LDA_test_acc=0.64
LDA_train_auc=0.739
LDA_test_auc=0.703
LDA_train_recall=0.56
LDA_test_recall=0.56
LDA_train_precision=0.69
LDA_test_precision=0.60
LDA_train_f1=0.61
LDA_test_f1=0.58
```

In [144]: index=[ 'Accuracy', 'AUC', 'Recall', 'Precision', 'F1 Score' ]

```
data = pd.DataFrame({'LDA Train':[LDA_train_acc,LDA_train_auc,LDA_train_recall,
                                LDA_train_precision,LDA_train_f1],
                    'LDA Test':[LDA_test_acc,LDA_test_auc,LDA_test_recall,LDA_test_precision,LDA_test_f1]})
round(data,2)
```

Out[144]:

|                  | LDA Train | LDA Test |
|------------------|-----------|----------|
| <b>Accuracy</b>  | 0.68      | 0.64     |
| <b>AUC</b>       | 0.74      | 0.70     |
| <b>Recall</b>    | 0.56      | 0.56     |
| <b>Precision</b> | 0.69      | 0.60     |
| <b>F1 Score</b>  | 0.61      | 0.58     |

In [ ]: